

archiving files and directories

- An archive file captures metadata and contents of multiple files and directories as a single file
- often incorporates compression
- example file formats include:
 - **ar** - used for libraries of relocatable binaries (.o files)
 - **shar** - usually software packages, Unix-like systems
 - self-extracting shell-script!
 - **tar** - general purpose, Unix-like systems
 - **zip** - general purpose, many platforms, includes compression
 - **deb** - software packages, Debian-family Linux distributions

tar - archive/unarchive files and directories

```
# capture files in assignment directory tree  
# -c create an archive  
# -f archive filename  
# -z compress with gzip  
$ tar -zcf assignment.tar.gz assignment  
$ cp assignment.tar.gz /tmp  
$ cd /tmp  
# extract files from archive  
# -x create an archive  
# -v (verbose) - print filenames when extracting  
# -f archive filename  
$ tar -xvf assignment.tar.gz  
...
```

- xz/unxz (-J option to tar)
 - algorithm Lempel–Ziv–Markovchain
 - good level of compression
 - slow to compress but uncompression fast
- bzip2/bunzip2 (-j option to tar)
 - algorithm: Burrows–Wheeler algorithm
 - faster to compress than xz but compression level not as good
- gzip/gunzip (-z option to tar)
 - algorithm: DEFLATE
 - compression level not as good as bzip2
 - very widely available on Unix-like machines
 - used for HTTP compression

curl - interact with web-servers

- curl lets you interact from command line with web and other servers

```
# fetch a file
$ curl -O https://cgi.cse.unsw.edu.au/~cs2041/examples.zip
# get other info
$ curl -I https://unsw.edu.au
HTTP/1.1 200 OK
Server: Apache/2.4.34 (Red Hat) OpenSSL/1.0.1e-fips PHP/5.6.25
X-Powered-By: PHP/5.6.25
# send data to web server
$ curl -X PUT -H 'content-type: txt/plain' https://google.com
# send cookies to web server
$ curl -b 'id=42' https://google.com
....
```

- curl has many other options
- wget provide similar functionality
 - Andrew likes curl better than wget - but little difference for most tasks

ssh - encrypted remote login

- ssh was written by Finnish university student Tatu Ylönen in 1995
- quickly adopted as an internet standard

```
$ ssh-keygen
```

```
Generating public/private rsa key pair.
```

```
Enter file in which to save the key (/home/andrewt/.ssh/id_rsa):
```

```
# leaves private key in $HOME/.ssh/id_rsa
```

```
# leaves public key in $HOME/.ssh/id_rsa.pub
```

```
$ cat $HOME/.ssh/id_rsa.pub
```

```
ssh-rsa AAAAB3NzaC1yc2EAAAABIwAAAQEAXL+t
```

```
....
```

```
# keep this file secret
```

```
$cat $HOME/.ssh/id_rsa
```

```
-----BEGIN RSA PRIVATE KEY-----
```

```
....
```

- Add public key to `$HOME/.ssh/authorized_keys` to allow for access without password.
- Can also run commands remotely:

```
$ ssh z1234567@login.cse.unsw.edu.au ls -las
```

rsync - efficiently copies files & directories

rsync efficiently copies files & directories locally or between machines (using ssh)

```
# mirror a directory tree in your CSE account
# -a preserves metadata & copies recursively
# -P shows progress
$ rsync -aP assignment/ login.cse.unsw.edu.au:assignment_backup/
```

If you run rsync command again it will only copy files which have changed.

If only a part of large file changed, will copy only the change (delta).

Many options, see `man rsync`

rsync - example copying only changed file

```
# create a directory a containing 1000 files
$ mkdir a
$ for i in $(seq 1 1000);do echo $i >a/$i; done
# copy directory a to director b
$ rsync -a a/ b/
# change one file
$ echo hello andrew >a/42.txt
# rsync will only copy changed files
$ rsync -av a/ b/
sending incremental file list
./
42.txt

sent 1,347 bytes  received 38 bytes  2,770.00 bytes/sec
total size is 305  speedup is 0.22
```

rsync - example copying only changed part of large file

```
# create a 100mb file
$ dd if=/dev/random bs=1M count=100 of=100_mb_file
# takes 25 seconds to copy it to CSE (40Mbps NBN)
$ time rsync 100_mb_file login.cse.unsw.edu.au:100_mb_file
real    0m24.943s
# repeat the rsync without changing the file - very fast
$ time rsync 100_mb_file login.cse.unsw.edu.au:100_mb_file
real    0m0.782s
# change a few bytes of the file
$ echo hello andrew >>100_mb_file
# rsync still fast
$ time rsync 100_mb_file login.cse.unsw.edu.au:100_mb_file
real    0m0.846s
```


Tools for Managing Processes

- *process* is an instance of an executing program
- on Unix-like systems each process had a unique number (pid)
 - pid's are smallish non-negative integer
- **ps** ... show process information
- **kill** ... send a signal to a process
 - typically used to terminate a process
 - signal 9 always terminates process

```
# using kill to logout
```

```
$ ps
```

PID	TTY	TIME	CMD
344024	pts/2	00:00:00	bash
346137	pts/2	00:00:00	ps

```
$ kill -9 344024
```

Tools for Managing Processes

- **pgrep** ... print PIDs of processes matching selection criteria
 - **pkill** send a signal to processes matching selection criteria
- **killall** ... also send a signal to a process with particular names
 - less powerful but more widely available than **pkill**
- **top** ... real-time monitoring of running process
 - or more easy to use **htop**

```
# find processes with python in their name
```

```
$ pgrep python
```

```
10787
```

```
20060
```

```
25975
```

```
27475
```

```
# kill processes with 2 consecutive vowels (don't do this!)
```

```
$ pkill -9 '[aeiou][aeiou]'
```

```
# kill processes with 2 consecutive vowels (don't do this!)
```

```
$ pkill -9 '[aeiou][aeiou]'
```

```
# kill all programs named teams or zoom
```

```
$ killall zoom teams
```

Linux Filesystem Layout

- **/home** - home directories for users on the system
- **/bin** - important system programs (scripts and binaries)
- **/usr/** - system programs and associated files
 - **/usr/bin** system programs
 - **/usr/local/bin** custom installed local programs
 - **/usr/lib** - libraries (linked with programs)
 - **/usr/include** - header files for C programs
 - ...
- **/etc** - holds configuration for system programs
- **/opt** - multi-operating system packages sometimes installed here
- **/var** - system files that regularly change, e.g.: log files, database files.
- **/tmp** - directory for temporary files - files removed on reboot

Linux Filesystem Layout

- **/root** - home directory for root user
- **/boot** - files need to boot operating system
- **/dev** - pathnames for hardware devices.
- **/media** - mount-point for removable device
- **/proc** - special filesystem with information about processes
- **/sys** - special filesystem with information about system

/dev - directory for devices

- Unix devices are manipulated by special files (usually) in **/dev**
 - e.g a disk might appear as **/dev/sda**

```
$ ls -l /dev
```

```
...  
brw-rw---- 1 root disk      8,    0 May 21 08:38 sda  
brw-rw---- 1 root disk      8,    1 May 21 08:38 sda1  
...  
crw-rw-rw- 1 root root      1,    3 May 21 08:38 null  
...  
crw-rw-rw- 1 root root      1,    8 May 21 08:38 random  
...  
crw--w---- 1 root tty       4,    0 May 21 08:38 tty0  
...  
crw-rw-rw- 1 root root      1,    5 May 21 08:38 zero
```

- **c** indicates character device - read/write bytes
- **b** indicates block devices - read/write blocks

some interesting “virtual” devices

- **/dev/null**
 - writes do nothing (are ignored)
 - reads return nothing
- **/dev/zero**
 - writes also do nothing, but use /dev/null
 - reads return bytes containing 0
- **/dev/random /dev/urandom**

some interesting “virtual” devices

```
$ xxd /dev/null
```

```
$ xxd /dev/zero|head
```

```
00000000: 0000 0000 0000 0000 0000 0000 0000 0000 .....
00000010: 0000 0000 0000 0000 0000 0000 0000 0000 .....
00000020: 0000 0000 0000 0000 0000 0000 0000 0000 .....
00000030: 0000 0000 0000 0000 0000 0000 0000 0000 .....
00000040: 0000 0000 0000 0000 0000 0000 0000 0000 .....
00000050: 0000 0000 0000 0000 0000 0000 0000 0000 .....
00000060: 0000 0000 0000 0000 0000 0000 0000 0000 .....
00000070: 0000 0000 0000 0000 0000 0000 0000 0000 .....
00000080: 0000 0000 0000 0000 0000 0000 0000 0000 .....
00000090: 0000 0000 0000 0000 0000 0000 0000 0000 .....
```

```
$ xxd /dev/random|head
```

```
00000000: 16a0 e6c3 16ac d43d a258 de1f 9ba8 a24b .....=.X.....K
00000010: 5661 3488 a6e4 fa2f 4204 fb80 16ad e0ec Va4.... /B.....
00000020: 9eb2 e60b 7267 1250 1954 47a4 75bb 79bc ....rg.P.TG.u.y.
00000030: 410c 334a b38d 4801 550a 2c83 35a3 e30d A.3J..H.U.,.5...
00000040: d603 31b7 6062 6c54 3b6f 0e00 bd4a 765a ..1.`b1T;o...JvZ
00000050: cfe6 d39c 89ba 9a02 66cd 5044 f417 30da .....f.PD..0.
00000060: eba5 df10 223a b963 7ad4 160c ae06 7660 ....":.cz.....v`
00000070: a4c4 352e 0252 615c 8e17 7b30 2e48 6fb3 ..5..Ra\..{0.Ho.
00000080: f9d5 d4f7 2913 73a9 3042 07de dde5 3537 ....).s.0B....57
00000090: 686c fba5 f41a b66a 7b68 2d48 3077 c672 h1.....j{h-H0w.r
```

fdisk - manipulate Disk partitions

- Disks are usually separated into separate regions called partitions.
 - allows parts of disk to be used for different purposes
- **fdisk** is a simple program to view or change partitions
- widely-used MBR partition format being replaced by GPT format
- see **gdisk** for GPT partitions
- gparted graphical tool handles both MBR and GPT

```
$ sudo fdisk -l /dev/sdi
```

```
Disk /dev/sdi: 28.66 GiB, 30752636928 bytes, 60063744 sectors
```

```
Disk model: Ultra
```

```
Units: sectors of 1 * 512 = 512 bytes
```

```
Sector size (logical/physical): 512 bytes / 512 bytes
```

```
I/O size (minimum/optimal): 512 bytes / 512 bytes
```

```
Disklabel type: dos
```

```
Disk identifier: 0x00000000
```

Device	Boot	Start	End	Sectors	Size	Id	Type
/dev/sdi1		32	60063743	60063712	28.7G	c	W95 FAT32 (LBA)

Beware: dangerous operation - have backups!

- **ext4** - mostly widely used Linux file-system
- **ext2/ext3** - older versions of ext4 - limited with less features
- **btrfs** - copy-on-write filesystem with interesting features
- **zfs** - filesystem which can span disks with interesting features
- **ntfs** default Windows file-system - can be accessed from Linux
- **fat** - older Windows filesystem
 - widely used for removable devices such as SD cards and USB keys
- **nfs** - network file system used to provide remote access to file system
- **sshfs** - remote file system layered on ssh
 - you can use to access your CSE file at home

mkfs - create your own filesystem on a partition

```
# /dev/sdi is a 32Gb flash drive
$ ls -l /dev/sdi*
brw-rw---- 1 root disk 8, 96 Aug  4 12:47 /dev/sdi
brw-rw---- 1 root disk 8, 97 Aug  4 12:47 /dev/sdi1
$ mkfs /dev/sdi1
mke2fs 1.45.6 (20-Mar-2020)
Discarding device blocks: done
Creating filesystem with 262144 4k blocks and 100096 inodes
Filesystem UUID: 66028671-cece-47ff-804c-4a3b7f9f0ea5
Superblock backups stored on blocks:
    32768, 98304, 163840, 229376
Allocating group tables: done
Writing inode tables: done
Creating journal (8192 blocks): done
Writing superblocks and filesystem accounting information: done
```

- **Beware:** overwrites (destroys) contents of disk - potential huge data loss!

mount - mount a file-system

- mount makes a file-system available below a point in the file-system
 - mount point must be an existing directory
- umount reverses this.

```
$ mkdir /tmp/i
$ sudo mount /dev/sdb1 /tmp/i
$ ls -l /tmp/i
..
$ umount /tmp/i
```

- distributions typically include a helper program to mount & unmount removable devices.

fsck - repair a file-system

- Power failure or other unexpected events may leave a filesystem in inconsistent state.
- **fsck** (file system check) checks and repairs a file-system.

```
# copy a random byte into raw file-system to simulate corruption
$ sudo dd if=/dev/random of=/dev/sdi1 bs=1 seek=1024 count=1
# filesystem will not mount
$ sudo mount /dev/sdi1 /tmp/i
mount: /tmp/i: wrong fs type, bad option, bad superblock on /dev/sdi1, missing co
# repair file system
$ sudo fsck /dev/sdi1
e2fsck 1.46.4 (18-Aug-2021)
ext2fs_open2: The ext2 superblock is corrupt
fsck.ext2: Superblock invalid, trying backup blocks...
/dev/sdi1 was not cleanly unmounted, check forced.
Pass 1: Checking inodes, blocks, and sizes
...
# filesystem will now mount
$ sudo mount /dev/sdi1 /tmp/i
```

- File system should not be in use (unmounted)

/etc/fstab - filesystem configuration

- Configures file systems on device to be mounted when system starts.

```
$ cat /etc/fstab
# device    mount-point  fs-type  options
/dev/sda1   /            ext4     noatime,errors=remount-ro 1 1
/dev/sda2   none        swap     sw                        0 0
```

- Must include a root file-system on /
- Usually includes a swap device.

Often uses a unique label for devices because device names can change if hardware reconfigured, e.g. more disks added.

```
$ cat /etc/fstab
UUID=36bcd9b9-de07-4de0-82c6-509000029f0e / ext4 defaults 1 1
```

/etc/passwd - user “database”

User information in **/etc/passwd**

Password hashes in **/etc/shadow**

Every user has unique number: **uid**

```
$ sed 2q /etc/passwd
root:x:0:0:root:/root:/bin/bash
daemon:x:1:1:daemon:/usr/sbin:/usr/sbin/nologin
$ sudo sed 2q /etc/shadow
root:$6$YiSiP7Pehz8aoe...../:18379:0:99999:7:::
daemon:*:18362:0:99999:7:::
```

- better to not edit `/etc/passwd` & `/etc/shadow` directly
- instead use (distribution-specific) tools, e.g on Debian-like systems:
 - add users with **adduser**
 - remove users with **deluser**

/etc/group - group database

Group information in **/etc/group**

```
$ head /etc/group
```

```
root:x:0:
```

```
daemon:x:1:
```

```
bin:x:2:
```

```
sys:x:3:
```

```
adm:x:4:
```

```
tty:x:5:
```

- Each group has unique number: **gid**
- Better to not edit **/etc/group** directly
- instead use (distribution-specific) tools, e.g on Debian-like systems:
- Add users to groups with **adduser**
- Also **addgroup delgroup**

The root user

Many system actions require root (uid == 0)

su allows you to execute commands as root or any other user.

sudo allows command to be run as root

Use cautiously - easy to damage system with commands run as root.

Edit sudo config file /etc/sudoers with **visudo**

```
# Adding user to sudo group should allow them to run sudo  
$ adduser andrewt sudo
```


Linux Distributions

- A distribution packages the Linux kernel with many other programs.
- Hundreds of linux distributions
- distributions used by CSE students include:
 - Debian - classic distro, used for CSE systems, runs on many architectures
 - Ubuntu - popular distro, based on Debian
 - Mint - based on Ubuntu
 - Arch - lightweight rolling release
 - Alpine - lightweight distro used in containers

The Debian Linux Distribution

- One of the oldest Linux distribution (1993)
- widely used & available for many platforms.
 - 10 CPU architecture officially supported
 - 14+ more CPU architecture unofficially supported
- Stable - new release every 2 yrs.
- Many derivative distributions
- packages total 300+ million lines of code

- A packages contains files that make up an application
- And build scripts to install/remove application.
- May contain metadata for managing the package.
- Used to install new applications onto a system
- Debian uses the **.deb** format (as do Ubuntu, Mint)
 - other distributions use other formats, e.g **rpm** (Red Hat) and **pkg** (Arch)
- distributions have their own package management tools
- Debian has
 - **dpkg** - low level tool - you probably don't need to use directly
 - **apt** - command-line tool - you probably want to use this
 - in scripts use **apt-get** (same options)
 - **aptitude** & **synaptic** - high level GUI tools

apt - Debian Package Manager

```
# update database of packages available
$ sudo apt update
# install a package + dependencies
$ sudo apt install <packagename>
# install a downloaded package file
$ sudo apt install ./package.deb
# uninstall package
$ sudo apt remove <packagename>
# update all packages
$ sudo apt dist-upgrade
# search for a package
$ sudo apt search <packagename>
# search for a package
$ sudo apt search <packagename>
# print package metadata including dependencies
$ sudo apt info <packagename>
```

dpkg - Debian Package Tool

- **dpkg** can be used to install individual **.deb** files
- **dpkg** is a low level tool, you mostly want **apt**
- **dpkg** can only check dependencies
- **apt** search package repositories for package
 - also downloads dependencies

```
$ curl -LO https://github.com/COMP1511UNSW/dcc/releases/download/2.7.8/dcc_2.7.8_  
$ ls -l dcc_2.7.8_all.deb  
-rw-r--r-- 1 andrewt 60180 Aug  4 11:18 dcc_2.7.8_all.deb  
$ sudo apt install ./dcc_2.7.2_all.deb
```

dpkg - inspecting package metadata

```
$ apt info dcc
```

```
Package: dcc
```

```
Version: 2.7.8
```

```
Priority: optional
```

```
Section: devel
```

```
Maintainer: Andrew Taylor <andrewt@unsw.edu.au>
```

```
Installed-Size: 77.8 kB
```

```
Depends: python3 (>= 3.6), gdb (>= 7.12), clang (>= 7.0), valgrind (>= 1:3.13)
```

```
Homepage: https://github.com/COMP1511UNSW/dcc
```

```
Download-Size: 60.2 kB
```

```
APT-Manual-Installed: yes
```

```
APT-Sources: /tmp/2041/dcc_2.7.8_all.deb
```

```
Description: compiler for novice C programmers
```

dcc compiles C programs using clang and adds explanations suitable **for** novice programmers to compiler messages novice programmers are likely to encounter and not understand. dcc also adds code to the binary which detects run-time errors and print information likely to be helpful to novice programmers, including printing values of variable in lines used near where the run-time error occurred.

- Linux containers provided by a set of kernel features
- linux containers allow running a processes in a different “world”
- you can run a process with effectively:
 - different root directory
 - e.g. /usr/bin can look different to process
 - different namespace
 - e.g different process ids
 - e.g different view of network
 - different uid
 - specified resource limits
- allows a program to be run which needs different packages to those installed

- Docker popular set of tools for working with Linux containers
- uses a union file system in clever way
 - overlays images so base images can be reused
- can specify dependencies for a program and produce self-contained image
 - e.g. can specify program requires python-3.9
- can run image on any Linux, Windows or OSX system with Docker installed
 - independent of what software is installed on that platform
- **dockerhub** provides sharing similar to **github
- Docker great for many purposes, but heavyweight (requires daemon)
 - other container tools may be useful depending on needs
- good to experiment with over the term-break
 - containers hard to use at CSE
 - disk space issues if every student has containers
 - docker not available - requires root
 - udocker - only supports some containers
 - podman - doesn't work with NFS
 - easy to install on your own machine
 - install instructions: <https://docs.docker.com/get-docker/>

Running Shell Command Inside and Outside a Container

```
$ whoami
andrewt
$ pwd
/home/andrewt
$ grep andrewt /etc/passwd
andrewt:x:517:517:andrewt,,,:/home/andrewt:/bin/bash
$ docker run -it --rm debian:bookworm bash -l
root@fa3e8e5e1b38:/# whoami
root
root@fa3e8e5e1b38:/# pwd
/
root@fa3e8e5e1b38:/# find / | wc -l
57208
root@fa3e8e5e1b38:/# cat /etc/passwd
root:x:0:0:root:/root:/bin/ash
bin:x:1:1:bin:/bin:/sbin/nologin
daemon:x:2:2:daemon:/sbin:/sbin/nologin
...
root@fa3e8e5e1b38:/# grep andrewt /etc/passwd
```

Mounting A Directory Inside a Container

```
$ ls /home/andrewt|wc
      87      92     819
$ docker run -it --rm debian:bookworm bash -l
root@fa3e8e5e1b38:/# ls /home/andrewt|wc
      87      92     819
root@fa3e8e5e1b38:/# cd /home/andrewt
root@fa3e8e5e1b38:/# echo hello Andrew >hello.txt
root@fa3e8e5e1b38:/# exit
$ cat hello.txt
hello Andrew
```

Example Dockerfile

```
# use Debian bookworm as our base image
FROM debian:bookworm

# install packages and create some directories and a file
RUN \
    apt-get update &&\
    apt-get install -q -y apache2 curl &&\
    apt-get -q -y clean &&\
    mkdir -p /var/run/apache2 /var/log/apache2 /var/lock &&\
    echo hello Andrew >/var/www/html/hello.html

# specify default command
ENTRYPOINT \
    apachectl start &&\
    bash -l &&\
    apachectl stop
```

source code for Dockerfile

Example - Building and Running the Container

```
$ curl http://127.0.0.1/hello.html
curl: (7) Failed to connect to 127.0.0.1 port 80 after 0 ms: Connection refused
$ docker build -t my_apache .
$ docker run -it my_apache
root@544075dcfbf5:/# curl http://127.0.0.1/hello.html
hello Andrew
root@544075dcfbf5:/# exit
logout
$ curl http://127.0.0.1/hello.html
curl: (7) Failed to connect to 127.0.0.1 port 80 after 0 ms: Connection refused
$
```

Example - Making a Network Port visible Outside the Container

```
$ curl http://127.0.0.1:1080/hello.html
curl: (7) Failed to connect to 127.0.0.1 port 1080 after 0 ms: Connection refused
$ docker run -it -p 1080:80 my_apache
root@1699e0605ed6:/
```

In another window

```
$ curl http://127.0.0.1:1080/hello.html
hello Andrew
```

Example - Supplying Webpages from Outside the Container

```
$ mkdir /tmp/content
$ for i in $(seq 0 9);do echo $i > /tmp/content/$i.html; done
$ docker run -it -p 1080:80 -v /tmp/content:/var/www/html my_apache
root@1699e0605ed6:/
```

In another window

```
$ curl http://127.0.0.1:1080/5.html
hello Andrew
```