

# Relational Database Management Systems

---

- What is an RDBMS?
- RDBMSs in COMP3311
- PostgreSQL Architecture
- SQLite Architecture
- Using PostgreSQL in CSE
- Managing Databases
- Managing Tables
- Managing Tuples
- Table Definition Example
- Exercise: Creating/Populating Databases
- Managing Other DB Objects

## ❖ What is an RDBMS?

---

A **relational database management system** (RDBMS) is

- software designed to support large-scale data-intensive applications
- allowing high-level description of data (tables, constraints)
- with high-level access to the data (relational model, SQL)
- providing efficient storage and retrieval (disk/memory management)
- supporting multiple simultaneous users (privilege, protection)
- doing multiple simultaneous operations (transactions, concurrency)
- maintaining reliable access to the stored data (backup, recovery)

Note: databases provide **persistent** storage of information

## ❖ RDBMSs in COMP3311

---

### PostgreSQL


- full-featured, client-server DBMS, resource intensive
- applications communicate via server to DB
- can run distributed and replicated
- follows SQL standard closely, but not totally
- extra data types (e.g. JSON), multiple procedural languages

### SQLite

- full-featured, serverless DBMS, light user of resources
- intended to be embedded in applications
- follows SQL standard closely, but not totally
- no stored procedures, add functions by embedding in apps

## ❖ PostgreSQL Architecture


PostgreSQL's client-server architecture:



COMP3311 21T1 ♦ RDBMSs ♦ [3/12]

## ❖ SQLite Architecture

SQLite's serverless architecture:



COMP3311 21T1 ♦ RDBMSs ♦ [4/12]

## ❖ Using PostgreSQL in CSE

Using your PostgreSQL server in CSE (once installed):

- login to `grieg`, set up environment, start server
- use `psql`, etc. to manipulate databases
- stop server, log off `grieg`

```
wagner$ ssh YOU@grieg
grieg$ source /srvr/YOU/env
grieg$ pg start
grieg$ psql myDatabase
... do stuff with your database ...
grieg$ pg stop
grieg$ exit
```

Need to run the command `priv srvr` once before the above will work

## ❖ Using PostgreSQL in CSE (cont)

PostgreSQL files (helps to understand state of server)

- PostgreSQL environment settings ...  
/srvr/YOU/env
- PostgreSQL home directory ... /srvr/YOU/pgsql/
- under the home directory ...
  - postgresql.conf ... main configuration file
  - base/ ... subdirectories containing database files
  - postmaster.pid ... process ID of server process
  - .s.PGSQL.5432 ... socket for clients to connect to server
  - .s.PGSQL.5432.lock ... lock file for socket
  - Log ... log file to monitor server errors, etc.

## ❖ Managing Databases

---

Shell commands to create/remove databases:

- `createdb dbname` ... create a new totally empty database
- `dropdb dbname` ... remove *all* data associated with a DB

(If no *dbname* supplied, assumes a database called *YOU*)

Shell commands to dump/restore database contents:

- `pg_dump dbname > dumpfile`
- `psql dbname -f dumpfile`

(Database *dbname* is typically created just before restore)

Main SQL statements in *dumpfile*: `CREATE TABLE`, `ALTER TABLE`, `COPY`



## ❖ Managing Tables

---

SQL statements:

- `CREATE TABLE table ( Attributes+Constraints )`
- `ALTER TABLE table TableSchemaChanges`
- `DROP TABLE table(s) [ CASCADE ]`
- `TRUNCATE TABLE table(s) [ CASCADE ]`

(All conform to SQL standard, but all also have extensions)

DROP. . CASCADE also drops objects which depend on the table

- objects could be tuples or views, but *not* whole tables

TRUNCATE. . CASCADE truncates tables which refer to the table

## ❖ Managing Tuples

SQL statements:

- `INSERT INTO table ( Attrs ) VALUES Tuple(s)`
- `DELETE FROM table WHERE condition`
- `UPDATE table SET AttrValueChanges WHERE condition`

$Attrs = ( attr_1, \dots attr_n )$        $Tuple = ( val_1, \dots val_n )$

*AttrValueChanges* is a comma-separated list of:

- *attrname* = *expression*

Each list element assigns a new value to a given attribute.

## ❖ Table Definition Example

Make a table to hold student data:

```
CREATE TABLE Student (  
    zid      serial,  
    family   varchar(40),  
    given    varchar(40) NOT null,  
    d_o_b    date NOT NULL,  
    gender   char(1) check (gender in ('M','F')),  
    degree   integer,  
    PRIMARY KEY (zid),  
    FOREIGN KEY (degree) REFERENCES Degrees(did)  
);
```

`serial` is a special type which automaticall generates unique integer values

## ❖ Exercise: Creating/Populating Databases

---

Do the following:

- create a database called `ex1`
- create a table `T` with two integer fields `x` and `y`
- examine the catalog definition of table `T`
- use `insert` statements to load some tuples
- use `pg_dump` to make a copy of the database contents
- remove the `ex1` database, then restore it from the dump

## ❖ Managing Other DB Objects

Databases contain objects other than tables and tuples:

- views, functions, sequences, types, indexes, roles, ...

Most have SQL statements for:

- `CREATE` *ObjectType name* ...
- `DROP` *ObjectType name* ...

Views and functions also have available:

- `CREATE OR REPLACE` *ObjectType name* ...

See PostgreSQL documentation Section VI, Chapter I for SQL statement details.

Produced: 10 Feb 2021