

Optiver 



Advanced C++ Programming

University of NSW Guest Lecture





Why am I Here Today?

- Optiver is proud to sponsor Advanced C++ Programming at the University of New South Wales.
- At Optiver developers design, build and maintain a world-class automated trading platform, mostly in C++.
- This means:
 - Designing, developing, testing and deploying their own systems.
 - Choosing appropriate algorithms and data-structures.
 - Optimising their systems for low-latency.
 - Employing up-to-date, industry-best practice.
- This course supports these skills and provides a great foundation for working with Optiver
- Optiver donates \$2,250 in prizes for the best performers in the course.



Agenda



How we use C++ at Optiver



Features of C++20 in Depth

Optiver 


How we use C++ at Optiver





How we use C++ at Optiver?

- We've built a world-class trading system, mostly in C++
- *Information* flows to us from an exchange.
- Our auto-traders estimate prices and determine if we wish to *execute* any order operations – that is: place, amend and/or delete orders.
- If so, those order operations are sent to the exchange.
- Rinse and repeat!
- The system is composed of numerous microservices which can be broadly grouped into three categories:
 1. Pricing
 2. Execution
 3. Risk Management

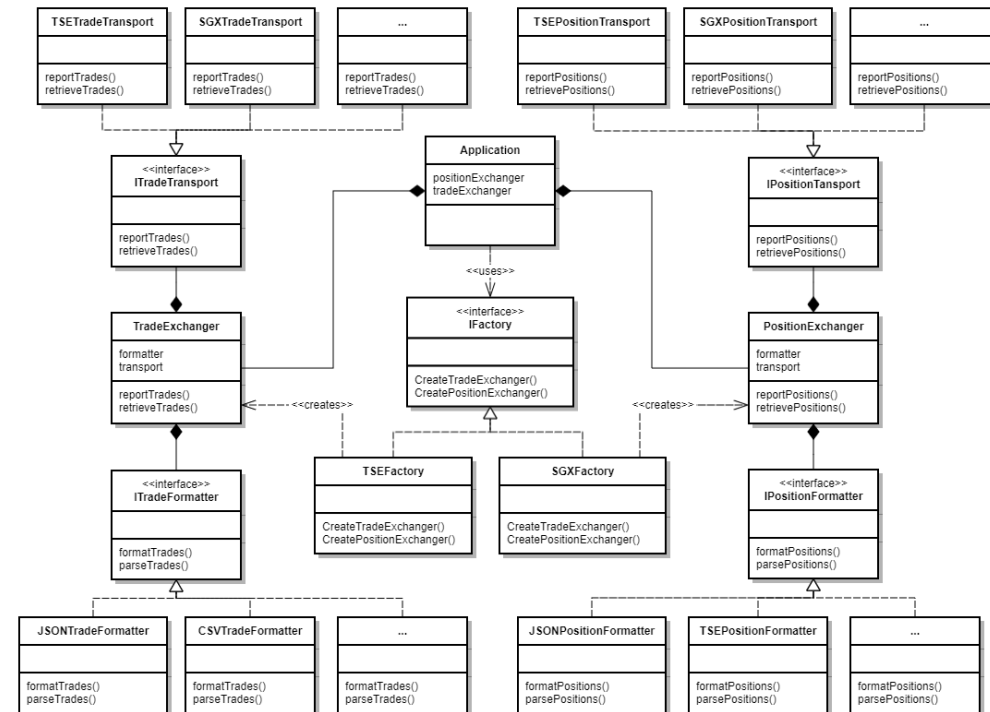




How we use C++ at Optiver

Object-Oriented Programming

- The features of C++ we use most are, of course, those that support object-oriented programming:
 - Classes
 - Encapsulation
 - Inheritance
 - Polymorphism
- We make heavy use of function overloading, and some use of operator overloading (e.g. the >> operator in logging)
- Object-oriented programming is considered critically important at Optiver and is specifically tested in our interview process





STL Containers

- Within our classes we make heavy use of STL containers such as `std::vector`
- The *latency* of our auto-traders is critically important and therefore we must choose containers very carefully

Let's look at an example:

- An important kind of information we deal with at Optiver is called an *order book*
- It describes the prices at which traders are willing to buy or sell and the volume they're willing to trade

| Buy Orders | Price | Sell Orders |
|------------|----------|-------------|
| | \$122.60 | 21 |
| | \$122.59 | 4 |
| | \$122.58 | 27 |
| | \$122.57 | |
| | \$122.56 | 10 |
| | \$122.55 | 3 |
| 15 | \$122.54 | |
| 10 | \$122.53 | |
| 20 | \$122.52 | |
| 8 | \$122.51 | |
| 27 | \$122.50 | |



An OrderBook Class


```
class OrderBook {
public:
    ...

    unsigned long get_volume(const std::string_view price_level) {
        auto it = mVolumes.find({price_level.data(), price_level.size()});
        if (it != mVolumes.end()) {
            return it->second;
        }
        return 0;
    }

private:
    std::map<std::string, unsigned long> mVolumes;
};
```



std::map – A Red-Black Tree





Use the Appropriate Kind of Map

```
class OrderBook {
public:
    ...

    unsigned long get_volume(const std::string_view price_level) {
        auto it = mVolumes.find({price_level.data(), price_level.size()});
        if (it != mVolumes.end()) {
            return it->second;
        }
        return 0;
    }

private:
    std::unordered_map<std::string, unsigned long> mVolumes; // <-- O(1) lookup
};
```



Be Aware of Hidden Costs

```
class OrderBook {
public:
    ...


    unsigned long get_volume(const std::string_view price_level) {
        auto it = mVolumes.find(price_level);    // <- No allocations
        if (it != mVolumes.end()) {
            return it->second;
        }
        return 0;
    }

private:
    std::vector<std::string> mPrices;
    std::unordered_map<std::string_view, unsigned long> mVolumes;
};
```




CPU Cache

- Modern CPUs employ caches to mitigate the cost of transferring data to/from memory.
- Data is transferred in blocks called cache lines (typically 64-bytes)
- Data structures in which all the data is located close together in memory are faster to access






`std::unordered_map` – A Hash Table





Circular Buffer



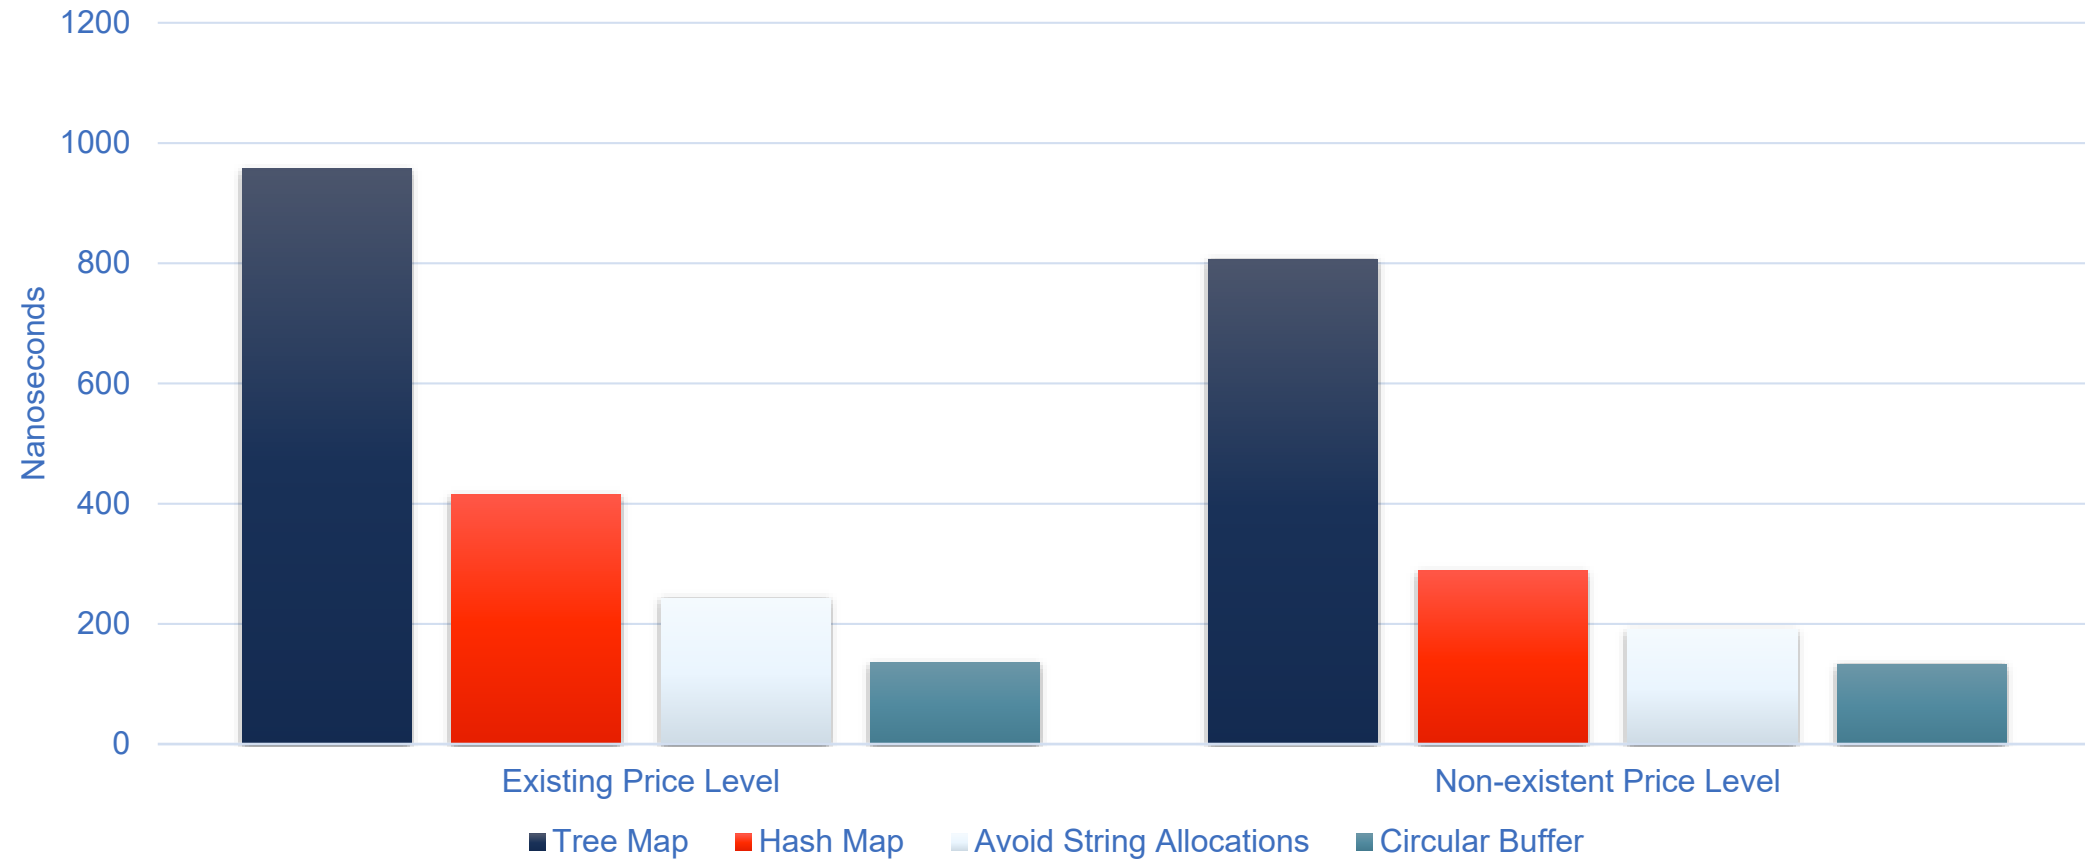


Take Advantage of the CPU Cache

```
class OrderBook {  
public:  
    ...  
  
    unsigned long get_volume(const std::string_view price_level) {  
        auto price = (unsigned long)(std::strtod(price_level.data(),  
                                                  nullptr)*100);  
        return mBuffer[price % mBuffer.size()];  
    }  
  
private:  
    std::vector<unsigned long> mBuffer;  
};
```




Benchmark Results of Order Book Variants





STL Algorithms

- We also make heavy use of STL algorithms such as `std::find`

To give a more complex example:

- One type of trading we do is called *base trading*
- We base our estimate of the fair price of one stock on the price of another (called the *base*)
- Calculating fair prices is time consuming, so we pre-compute them to save time, but the range of possible base prices is very large so we can't feasibly pre-compute them all
- Instead, we pre-compute fair prices for a small number of equidistant base prices that are near the current base price
- How can we quickly get a *theo* given a *base price*?

| Base Prices | | | | |
|-------------|---------|---------|---------|---------|
| \$42.25 | \$42.30 | \$42.35 | \$42.40 | \$42.45 |

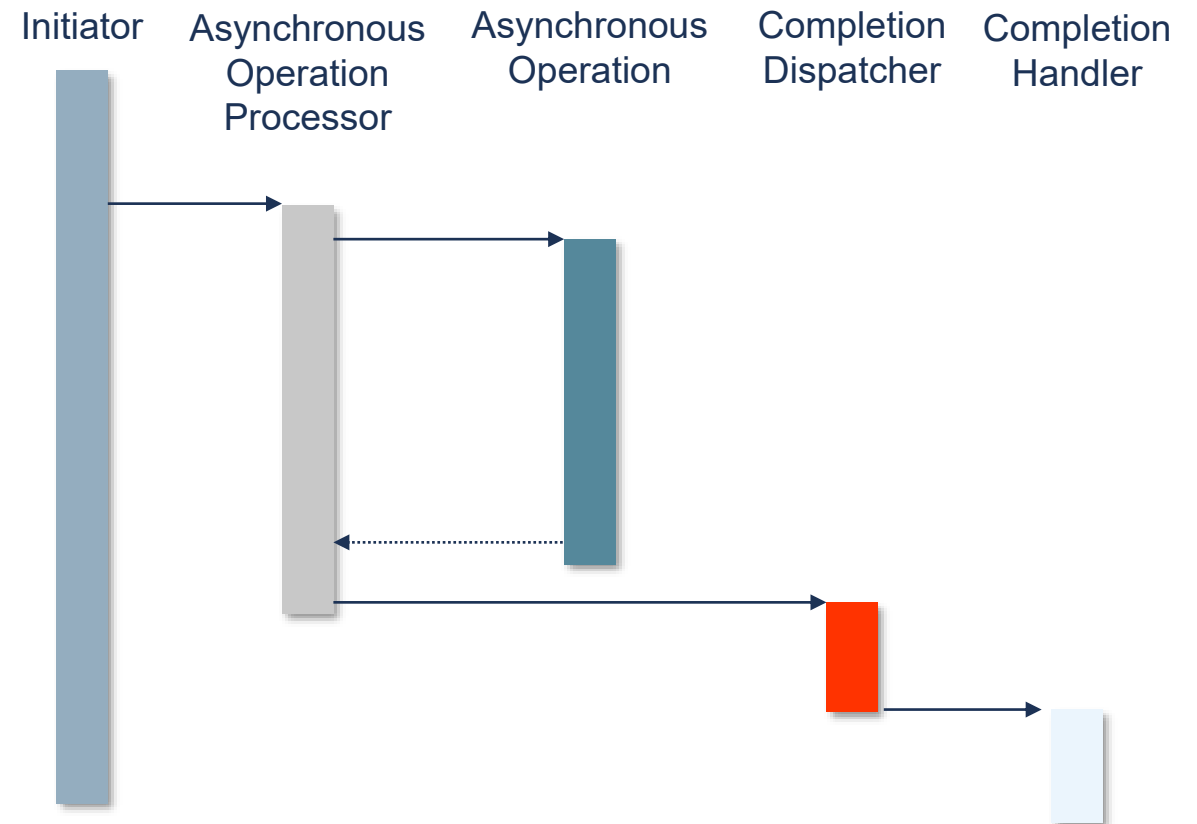
| Theoretical Prices | | | | |
|--------------------|-------|---------|---------|---------|
| \$63.80 | 63.84 | \$63.90 | \$63.96 | \$64.01 |



Lambda Functions

- One common way we use lambda functions is in the *Proactor Pattern*
- The *initiator* starts an asynchronous operation and defines a *completion handler* (often a lambda)
- The *completion handler* is called at the end of the operation
- One good example is when we wish to update a database:
 - A message is sent to the database requesting the update
 - A lambda function is defined which is called upon notification that the database update has completed successfully

The Proactor Pattern:





How we use C++ at Optiver

Other C++ Features Used Frequently

- Smart pointers, especially `std::unique_ptr`
 - Memory allocation can be time consuming so we tend to do most allocation at startup
- template metaprogramming
 - e.g., to define a database ORM
- `auto`

Optiver 

C++20 Features in Depth





Modules

- Standardised mechanism for code reuse
- Organise C++ code into logical components
- A module explicitly exports the classes and functions that code outside of the module are allowed to access.
- Other code remains private
- This behaviour will be extended to all c++ library header files (pending compiler support)

```
// helloworld.cppm
export module helloworld; // module declaration
import <iostream>;        // import declaration

export void hello() {      // export declaration
    std::cout << "Hello world!\n";
}

void goodbye() {
    std::cout << "Goodbye world!\n";
}

//main.cpp
import helloworld; // import declaration

int main() {
    hello();
    goodbye(); //error not exported
}
```



Modules

- Reduces dependency on the pre-processor and header files
- #include file order no longer matters, less error prone, no cyclic dependencies
- Avoids issues with macro leaking
- Faster build times, modules are pre-compiled only once

```
// helloworld.cppm
export module helloworld; // module declaration
import <iostream>;        // import declaration

export void hello() {     // export declaration
    std::cout << "Hello world!\n";
}

void goodbye() {
    std::cout << "Goodbye world!\n";
}

//main.cpp
import helloworld; // import declaration

int main() {
    hello();
    goodbye(); //error not exported
}
```



Coroutines

- A coroutine is a function that can suspend execution to be resumed later
- Control is returned to the caller
- The current state of the coroutine is saved to be resumed where it left off
- Keywords `co_yield`, `co_await`, `co_return`

```
generator<int> iota(int n = 0)
{
    while (true)
    {
        co_yield n++;
    }
}

void printNumbers()
{
    for (const int i : iota())
        std::cout << i << std::endl;
}
```




Coroutines

- Stackless - Coroutine invocations do not have independent stacks, they allocate data for the coroutine on the heap – efficient memory usage and context switching
- Allow for sequential code that executes asynchronously
- No callbacks, can yield control and resume when necessary

However:

- C++ 20 only provides a very low-level api, the generator class used here is from C++23
- Rules of interaction between the caller and the callee are complex
- Can use some third party libraries (e.g. cppcoro)

```
generator<int> iota(int n = 0)
{
    while (true)
    {
        co_yield n++;
    }
}

void printNumbers()
{
    for (const int i : iota())
        std::cout << i << std::endl;
}
```



Coroutines

- At Optiver, we make heavy use of callbacks to handle events
- Coroutines allow us to declare business logic in-line rather than split across multiple callbacks
- Coroutines have application defined context switch points, which gives us concurrency with a single thread

```
task<> tcp_echo_server() {  
    char data[1024];  
    while (true)  
    {  
        std::size_t n = co_await  
            socket.async_read_some(buffer(data));  
        co_await async_write(  
            socket, buffer(data, n));  
    }  
}
```



Concepts

- Concepts allow us to specify what is needed from a template argument so this can be checked by the compiler
- Constraints model semantic requirements
- In this example the parameter *T* is unconstrained, but it won't compile for any type that doesn't have a + operator
- These error messages can be very complex

```
template <typename T>
auto add(T const a, T const b)
{
    return a + b;
}

int main()
{
    std::cout << add(1, 3) << std::endl;
}
```



Concepts

- Add a *requires* clause

```
template <typename T>
requires std::integral<T>
auto add(T const a, T const b)
{
    return a + b;
}

int main()
{
    std::cout << add(1, 3) << std::endl;
}
```



Concepts

- Creating our own concept
- Generates meaningful error messages that are much easier to understand
- Clearly documents expectations
- Much easier to use than previous **enable_if** syntax

```
template <typename T>  
concept Number = std::integral<T> || std::floating_point<T>;
```

```
template <typename T, typename U>  
requires Number<T> && Number<U>  
auto add(T const a, U const b)  
{  
    return a + b;  
}  
  
int main()  
{  
    std::cout << add(1, 3.1) << std::endl;  
}
```



Ranges

- Ranges are an abstraction of a "collection of items" or "something iterable"
- Containers are ranges, they own their elements
- Views are ranges that are usually defined on another range
- Views do not own any data beyond their algorithm
- Less error prone than using iterators

```
std::vector v;  
std::sort(v.begin(), v.end());  
std::ranges::sort(v);
```



Ranges

- Allow us to lazily filter and transform data through a pipeline
- Views are applied when an element is requested, not when the view is created

```
#include <iostream>
#include <ranges>
#include <vector>

int main() {
    std::vector<int> numbers = { 1, 2, 3, 4, 5, 6 };

    auto is_even = [](int n) { return n % 2 == 0; };

    auto results = numbers | std::views::filter(is_even)
        | std::views::transform([](int n) { return n++; })
        | std::views::reverse;

    for (auto v: results) {
        std::cout << v << " ";    // Output: 7 3 5
    }
}
```



Questions?




Greg Saunders
Head of IT Education


Greg.Saunders@optiver.com.au



www.optiver.com/working-at-optiver









| <i>Table header first column in italic (manually)</i> | Header row | (right aligned) | Last column |
|---|------------|-----------------|-------------|
| | | | |
| | | | |
| | | | |
| | | | |
| Subtotal bold and lightblue hatching must be set manually | | | |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |
| Total row | Total row | Total row | Total row |



| <i>Table header first column in italic (manually)</i> | Header row | (right aligned) | Last column |
|---|------------|-----------------|-------------|
| | | | |
| | | | |
| | | | |
| | | | |
| Subtotal bold and lightblue hatching must be set manually | | | |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |
| Total row | Total row | Total row | Total row |

FOOTNOTES
Text



Do you have
any questions?

Name Surname
Job title

namesurname@optiver.com
+31 (0) 625755624
Strawinskylaan 3095
1077 ZX Amsterdam